

# Diffusion Models for Generative AI

Taif Alsulami<sup>1</sup>, Rahaf Banat<sup>2</sup>

CS3012: Algorithm Analysis

Department of Computer Science College of Engineering Effat University, Fall 2024

Supervisor: Dr. Mohammad Amjad

**Abstract**—Diffusion Models are a new, powerful approach to generative artificial intelligence that has been dealt with in the creation of high-quality images, videos, and other contents using a probabilistic framework that systematically adds and removes noise. This paper attempts to delve into the utility, novelty, and intricate workings of diffusion models and their immense importance across a wide spectrum of tasks and fields ranging from content creation to scientific simulations. We will look through the forward diffusion and reverse denoising processes, architectural components, and diffusion models in comparison to other generative models like GANs, in order to bring out unique advantages and applications that diffusion models have in modern AI.

**keywords:** Generative AI, Diffusion Models, Image Synthesis, Noise-Based Generation, Denoising Process, Probabilistic Modeling, Content Creation, GAN Comparison, AI Applications

## I. INTRODUCTION

Generative AI has revolutionized data synthesis, where machines can now produce on their own content that closely emulates human-generated data sometimes to a degree exceeding human expectations. The traditional frontrunners in this field, GANs and VAEs, have helped in many ways and supported progress in image synthesis, video generation, and more. For example, GANs use a generator-discriminator setup that, through optimization, allows the generation of photorealistic images based on a game-like interaction between the two components. Nevertheless, some challenges arise with GANs: their dependence on adversarial training might make them unstable to train and, in most cases, very delicate to balance, which often collapses into mode collapse, where the generator always produces very few varieties of images. This might compromise the diversity of generated outputs, rendering GANs less reliable in tasks requiring high variability. On the other hand, VAEs approach data generation via a probabilistic framework; hence, in theory, allowing smooth interpolations between the generated data points is possible. However, the VAEs are not normally good at generating high-resolution images since their structure does not always capture fine-grained details in complex data distributions.

Diffusion models have recently been brought forth as a powerful alternative in the arena of generative AI. Their noise-based, probabilistic framework takes the input data and adds noise to it, then progressively removes the noise to achieve the transformation of data. Originally proposed by Sohl-Dickstein, diffusion models were not high in the spotlight until very recently, when new computational tools started to make possible the practical application of such models. Unlike

GANs, which rely on adversarial training, diffusion models use a structured forward and reverse process to stabilize training and enable generating high-quality synthetic data. The current paper will present a comprehensive survey of diffusion models and discuss architectural innovations, unique benefits, and the increasing number of applications demonstrating their relevance in generative AI. We bring forth their stability and high fidelity through analysis of diffusion models together with traditional generative models, hence showing potential as a foundational technology in AI-based content generation, scientific modeling, and beyond. [1]

## II. HISTORY OF DIFFUSION MODELS

Diffusion models draw on the study of stochastic processes. Initially, in fields such as physics and finance, diffusion became an important tool for modeling random phenomena, starting with the Brownian motion of particles and going up to the behavior of returns in security markets. Foundational ideas had their contribution in early AI research; however, it was only till 2015 that a "diffusion" process for generative modeling has been proposed by Sohl-Dickstein. The approach involved a process: slowly corrupting structured data into noise while learning a reverse process to regenerate the data from the noise. This model was a structured and stable alternative to GANs and VAEs, but computationally intensive and complex; thus, it was not quite immediately practical and, for quite some time thereafter, largely ignored. [2]

Interest in diffusion models suddenly grew up until around 2020, catalyzed by improvements in computational power and model optimization. Jonathan Ho developed the original approach into "Denoising Diffusion Probabilistic Models" (DDPM), significantly more efficient in the reverse process and generating high-quality images that, for the first time, could compare to GANs regarding fidelity and stability. This integration with U-Net architectures and attention mechanisms has further pushed the ability of diffusion models on high-dimensional data management, making them applied to subdomains from image synthesis to video generation and biomedical imaging. Diffusion models have made their way to the fore through successful applications such as DALL-E and Stable Diffusion for their versatility and high fidelity, marking them as a core generative AI technology. [3]

### III. BACKGROUND AND RELATED WORK

The advancement of generative models has reached several key milestones. GANs gained popularity for image generation thanks to their capability to create high-quality, realistic images. However, GANs require precise balancing during their adversarial training to avoid problems such as mode collapse, where the generator creates a narrow range of images [4]. Variational Autoencoders (VAEs) offered a probabilistic method for generating data; nonetheless, they faced challenges in accurately capturing the complex nuances of intricate datasets due to their inherent structural constraints.

In contrast, diffusion models conceptualize data generation as a gradual transformation from noise into structured data. This gradual diffusion process improves stability and eliminates the need for adversarial training. Recent advancements have enhanced diffusion models, increasing their sampling efficiency and expanding their use across multiple fields [5].

In contrast, diffusion models represent a shift from adversarial and probabilistic generation to a noise-based framework. Rather than producing data outright, diffusion models view the process as a slow evolution from random noise to organized data. This structured diffusion methodology offers remarkable stability, negating the necessity for adversarial training and sidestepping challenges such as mode collapse. Recent research has improved diffusion models, increasing their sampling efficiency and broadening their use across various sectors. In contemporary studies, diffusion models have attracted considerable interest for their roles in biomedical imaging, scientific simulations, and creative content creation, where high fidelity and dependability are essential.

### IV. ALGORITHMIC PROCESS OF DIFFUSION MODELS

Diffusion models are based on a well-defined algorithmic process, which can be conceptualized as two major phases: the forward diffusion process and the reverse denoising process. In the forward process, it adds noise in steps to structured data until it gets transformed into pure noise. Most commonly, it is represented as a Markov chain, systematically breaking down data structure over multiple timesteps following a probabilistic model. At each step  $t$ , a small amount of Gaussian noise is added to the data, in a process that incrementally reduces its identifiable features until it eventually resembles random noise. This degradation path provides a "blueprint" that the model will later use to reconstruct the data in reverse. Mathematically, the forward algorithm can be expressed as a sequence  $x_0, x_1, \dots, x_T$ , where each  $x_t$  is a noisy version of the original data point  $x_0$ . By performing the iterations over these steps, the model constructs a structured noise path that will form the basis for its reverse denoising algorithm.

The reverse process is where the generative power of the model comes into action. Starting from a noise sample, the reverse process would progressively eliminate the noise and build up the original data by following the path that is strictly opposite to the forward diffusion process. This process is

guided by a neural network that has been trained to estimate the level of noise in every step in time, removing it step by step until all noise has been removed and data restored. This reverse process is actually very expensive to compute in practice, since each of the denoising steps in the reverse process should be accompanied by the accurate estimation of the noise added in the forward process to stay faithful to the data distribution. Recent algorithmic improvements, including score-based denoising and faster sampling, have been helpful in making that more viable in practice. Both the forward and reverse processes constitute a single algorithmic framework, whereby the diffusion models synthesize high-quality and realistic data by a sequence of probabilistic transformations. [6]

### V. WORKING OF DIFFUSION MODELS

#### A. Forward Diffusion Process

In the forward diffusion process, the initial data experiences a progressive decline due to a controlled increment of Gaussian noise at every time step. This addition of noise changes the data into a form that is indistinguishable from random noise. The forward process enables the model to approach the data generation task by converting random noise into organized data. This progressive degradation is mathematically expressed by the transition  $x_0 \rightarrow x_t$ , where  $x_t$  represents the data at time step  $t$ , and noise is applied incrementally.

The model gains the transition dynamics needed to faithfully reconstruct the data in the reverse phase by converting the initial data into noise. The forward diffusion process lays the groundwork for comprehending the data's structure, successfully creating a noise pathway that the model will reverse during generation [7].

$$q(x_t|x_{t-1}) = \mathcal{N}(x_t; \sqrt{\alpha_t}x_{t-1}, (1 - \alpha_t)I) \quad (1)$$

#### B. Reverse Denoising Process

The generative phase, the reverse denoising process, consists of the model generating data by reversing the noise added during the forward phase. Beginning with a sample of noise, the model gradually removes noise at each time step, restoring the original data. Although this iterative noise reduction is computationally intensive, it provides precise control over the generated outputs, resulting in high-quality results. The denoising process is influenced by probabilistic principles, with each step moving the model closer to the original data distribution.

For the reverse process to work effectively, the model must accurately assess the noise added at each time step, guaranteeing that the reconstructed data closely matches the original input. This denoising method gives the model a robust framework for generating intricate and lifelike outputs across various data types [8].

$$p_\theta(x_{t-1}|x_t) = \mathcal{N}(x_{t-1}; \mu_\theta(x_t, t), \Sigma_\theta(x_t, t)) \quad (2)$$




Fig. 1. Diffusion Process: Adding and Removing Noise "Illustrating the forward diffusion and reverse denoising processes




Fig. 2. diffusion model

### C. Model Architecture

Diffusion models often utilize a U-Net architecture combined with attention mechanisms. Initially created for segmenting biomedical images, the U-Net design is adept at maintaining fine details while ensuring a broader context, which makes it especially suitable for generation based on noise. The encoder-decoder design of the model captures features across different scales, ensuring that the data's local and global characteristics are preserved during the diffusion process.

Attention mechanisms also improve the model's capacity to concentrate on pertinent data sections. By emphasizing specific areas, these mechanisms enable the model to produce unified images with intricate details, making it highly effective for handling high-dimensional data such as images [9]. By integrating U-Net and attention, diffusion models achieve an architecture that is adept at handling complex, high-dimensional data.

## VI. NOVELTY AND UTILITY OF DIFFUSION MODELS

### A. High Fidelity and Stability

The structured approach diffusion models use ensures systematic and uniform training, preventing the adversarial issues often found in GANs. By following a specified noise pathway, these models yield high-quality outcomes, making them appropriate for essential uses that demand precision and dependability, like medical imaging and content creation [10].

### B. Versatile Applications

Diffusion models find applications across multiple domains, such as the creative arts, healthcare, and scientific research. They allow artists to produce lifelike images and intricate animations within the creative industry. In the healthcare sector, these models assist in generating synthetic data for training machine learning algorithms, which is beneficial for tasks like diagnostic imaging while maintaining patient confidentiality. Their versatility makes diffusion models important across different fields, enabling high-quality synthesis where precision is critical [11].

## VII. COMPLEXITY AND ANALYSIS

### A. Computational Complexity

The process in the time-consuming denoising step of the diffusion models has a corresponding reverse step that requires an accurate estimate with noise reduction. Although this process is iterative to meet high-quality standards, it has been quite expensive in terms of computation. More recent works have concentrated on increasing sample efficiency by simplifying model designs to reduce computational processing burdens. While other powerful diffusion models have become possible because of strategies such as faster sampling algorithms and simplification in network designs, further work is necessary to make them practical for real-time applications. [12].

### B. Comparison with GANs and VAEs

Diffusion models have clear advantages over GANs and VAEs: they avoid the problems inherent in using an adversarial setup by instead relying on a structured noise-based method. This can cause GANs to suffer from mode collapse and requires a very well-balanced training to be performed. Diffusion models are less vulnerable to the dangers of adversarial training because of this stability. Despite coming from probabilistic modeling, VAEs' low resolution prevents them from producing high-quality pictures. Although they are more expensive to compute, diffusion models are superior to GANs and VAEs because they employ a step-by-step process based on sounds, which makes them more realistic and detailed. [13]




Fig. 3. Overview of different types of generative models

## VIII. APPLICATIONS OF DIFFUSION MODELS

### A. Media and Art Generation

Diffusion models have quite literally reshaped the content creation landscape in the creative industries by enabling the production of high-resolution, photorealistic images and videos. Platforms like DALL-E offer a service that enables artists, marketers, and all types of creators to use diffusion models to generate images from simple text prompts. This functionality expands the possibilities for artistic projects by making it possible to generate unique graphics while requiring less human labor. [14]

### B. Healthcare and Biomedical Imaging

In the field of healthcare, diffusion models provide privacy-preserving synthetic data generation for training machine learning models in diagnostic imaging. The diffusion models allow generation of synthetic scans, which can be used by medical institutions to enhance diagnostic algorithms without letting go of patient confidentiality. High-fidelity diffusion models provide these synthetic images, fit for use in medical applications where precision and accuracy are a necessity. [15]

### C. Scientific Research and Simulation

Diffusion models are obviously of importance in scientific inquiry and, in data synthesis for complex simulations. In studies dealing with molecular biology and climate science, these models help scientists to create data reflecting reality—complex natural processes that enhance visualization in order to foster deeper understandings of complex systems. The ability to produce precise depictions of natural phenomena underscores the role of diffusion models as an essential instrument for scientific inquiry. [16]

## IX. ETHICAL CONSIDERATIONS

Ethical considerations include privacy, responsible AI practices, and the potential for misuse, which are critical when deploying generative models in real-world applications. The paper addresses these ethical implications, especially in sensitive healthcare areas where patient data and privacy must be rigorously protected. It is crucial to ensure transparency and for users to understand the strengths and limitations of generative models to avoid unforeseen outcomes. In addition, the document suggests a series of ethical guidelines aimed at preserving user trust, promoting accountability, and adhering to best practices in the responsible development of AI, thereby reducing risks linked to bias and data privacy. [17]

## X. ADVANTAGES OF DIFFUSION MODELS

Diffusion models provide several benefits that increase their attractiveness for generative tasks in artificial intelligence. A key benefit is their **stability in training**. Unlike GANs, which depend on a complex adversarial training method that may result in instability and mode collapse, diffusion models utilize a systematic, non-adversarial framework. This stability enables diffusion models to consistently produce high-quality outcomes without the potential of narrowing down to a limited

range of options, making them more reliable for a range of applications.

Another important advantage is the **high fidelity** of the outputs produced by diffusion models. The stepwise denoising method provides exact control over the produced data, enabling diffusion models to capture global structures and intricate details effectively. This quality is invaluable for applications that demand high-resolution outcomes, such as medical imaging, video creation, and the production of artistic work.

Diffusion models also shine in terms of **data diversity** and **flexibility in generation**. They can create various outputs by employing probabilistic pathways and conditioning methods while staying true to the data distribution without being confined to specific modes or patterns. The versatility of diffusion models allows them to be utilized across multiple sectors, including biomedical imaging, scientific simulations, and creative industries. Their systematic method facilitates the smooth incorporation of various data types, including text, audio, and video, thereby expanding their possible applications.

Additionally, recent progress in **algorithm optimization** has enhanced the computational efficiency of diffusion models, making them more applicable to real-world scenarios. Methods like variance reduction, score-based denoising, and accelerated sampling algorithms have lessened the time and resources needed for these models to produce high-quality results. This advancement increases their practicality and creates opportunities for real-time applications, where both speed and quality are crucial. [18]

## XI. DISADVANTAGES OF DIFFUSION MODELS

Despite their many strengths, diffusion models also have disadvantages that limit their broader adoption and application. One of the most notable drawbacks is their **high computational cost**. The reverse denoising process's iterative nature requires many steps, each involving complex calculations to remove noise accurately. This computational demand makes diffusion models resource-intensive, often requiring powerful hardware setups or extensive cloud resources, which can be cost-prohibitive for some applications.

One drawback is the **time complexity** linked to sampling in diffusion models. The necessity to progressively reverse the noise addition process results in prolonged generation times, particularly when contrasted with GANs, which can produce samples in a single forward pass after training. In real-time applications like interactive content creation or on-device AI, the slower generation rates of diffusion models can pose a considerable drawback, making them less ideal for situations that demand immediate feedback or quick generation.

AdFurthermore, **interpretability** is a challenge associated with diffusion models. These models' probabilistic and incremental nature can make understanding the specific mechanisms involved in the generation process challenging. In contrast to simpler generative models, where specific components or parameters may have a more apparent influence on the output, diffusion models entail intricate transformations

over several stages, making their analysis and debugging more complex.

Ultimately, although diffusion models offer a wide range of applications, they often necessitate **fine-tuning and a substantial amount of data** to reach the best outcomes in particular scenarios. This reliance on data can pose a challenge in areas where it is limited or hard to acquire, as the model's efficacy depends on having an extensive and varied dataset. When provided with adequate data, diffusion models can grasp the intricacies of producing high-quality results, especially in specialized domains such as medical imaging or scientific simulations. [19]

## XII. FUTURE DIRECTIONS IN DIFFUSION MODELS AND GENERATIVE AI

Current research is now focused on improving the efficiency of diffusion models, increasing their scope, and opening up new avenues within generative AI. In real-time generation, the focus of the next wave of development will ensure that diffusion models become applicable in interactive domains such as live video editing, immersive virtual reality, and interactive content creation. Next would come multi-modality diffusion models that would allow the creation of synchronized data across formats, from text to image or audio to video. This could result in frictionless, high-quality, multimodal outputs, enabling new experiences in areas like education and entertainment, to AI-powered personal assistants reacting with rich, cross-modal responses.

Another exciting frontier in the domain involves neural architecture searching for diffusion models. The NAS automatically further develops the design of optimal architecture and might discover highly efficient structures that further reduce computational demands. In addition, with increasing resource efficiency, there is a possibility for the deployment of diffusion models on edge devices, such as mobile phones or IoT systems, for application. This will make users able to access high-quality generative models without dependence on powerful central servers, thereby democratizing high-level AI capabilities. [20]

## XIII. DIFFERENCES BETWEEN DIFFUSION MODELS AND OTHER GENERATIVE MODELS

Diffusion models differ fundamentally from GANs and VAEs in their generation techniques and training methodologies. In contrast to GANs, which function within an adversarial setup featuring a generator and a discriminator, diffusion models utilize a systematic approach based on noise that gradually adds and removes noise to produce realistic samples. This method alleviates the instability and mode collapse problems frequently encountered with GANs. Furthermore, it improves the output fidelity since these models do not depend on a traditional latent space like VAEs, which may lead to loss of detail during the encoding and decoding processes.

Diffusion models offer a more stable training alternative than GANs, which struggle with balancing generator and discriminator networks, often leading to mode collapse where

output diversity suffers. By using a probabilistic framework without a discriminator, diffusion models achieve reliable training dynamics but at the cost of computational efficiency due to their iterative denoising process. Simultaneously, once GANs are trained, they can generate samples in a single forward pass. GANs may show limited diversity even with high-quality synthesis regarding output variety and control. Although VAEs facilitate the manipulation of latent variables for smooth transitions, they often produce images that lack detail. Diffusion models are outstanding in creating both diverse and highly intricate outputs by following a systematic noise process. They can integrate conditioning mechanisms, such as those utilized in text-to-image generation, which enhances control and adaptability across different applications. [21]

## XIV. CONSTRAINTS AND LIMITATIONS

Acknowledging constraints and limitations in computing power and resource availability affects the practicality of diffusion models for smaller-scale applications. Typically, these models necessitate significant memory, processing capabilities, and storage to effectively manage the complex procedures of generating high-quality content. This section examines the effects of these limitations, particularly for applications running on resource-limited devices like mobile or edge platforms. Alternative strategies for resource-constrained environments are explored, including optimization methods such as model quantization and pruning to lessen computational requirements while preserving the model's effectiveness. Furthermore, hybrid and cloud-based solutions are evaluated, allowing smaller devices to utilize sophisticated generative models without placing direct processing burdens on them.

## XV. CONCLUSION

Diffusion models are a potential generative AI technology. They allow for a stable, high-quality way to generate data through structured noise-based processes. These models circumvent the problems that GANs are beset with, and they perform better than VAEs at capturing details. This makes diffusion models quite powerful in generating realistic, high-quality results. Such versatility across media, healthcare, and scientific fields points out their potential as a transformation tool in AI-driven content creation, simulation, and data synthesis. Future research is likely to focus on increasing the efficiency of the diffusion models so that they can be made accessible for real-world, real-time applications and extended in applicability toward an even broader range of domains. Diffusion models are improving and are probably going to be a huge deal in generative AI.

## REFERENCES

- [1] A. Smith, B. Johnson, and C. Lee, "Diffusion Models in Generative AI," *ResearchGate*, 2024. [Online]. Available: [https://www.researchgate.net/publication/375030647\\_Diffusion\\_Models\\_in\\_Generative\\_AI](https://www.researchgate.net/publication/375030647_Diffusion_Models_in_Generative_AI).
- [2] . Vincent, "A Brief History of Diffusion, the Tech at the Heart of Modern Image-Generating AI," *TechCrunch*, Dec. 22, 2022. Available: <https://techcrunch.com/2022/12/22/a-brief-history-of-diffusion-the-tech-at-the-heart-of-modern-image-generating-ai/>.

- [3] IBM, "Diffusion Models: The Next Step in AI," *IBM Think Blog*, 2024. Available: <https://www.ibm.com/think/topics/diffusion-models>.
- [4] L. Yang, Z. Zhang, Y. Song, et al., "Diffusion Models: A Comprehensive Survey of Methods and Applications," *ACM Computing Surveys*, 2024.
- [5] H. Cao, C. Tan, Z. Gao, et al., "A Survey on Generative Diffusion Models," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 2023.
- [6] Y. Zhang, B. Hartl, H. Hazan, and M. Levin, "Diffusion Models are Evolutionary Algorithms," *arXiv preprint arXiv:2410.02543*, Oct. 2024. Available: <https://arxiv.org/abs/2410.02543>
- [7] MIT News, "AI Generates High-Quality Images 30 Times Faster in a Single Step," MIT News, Mar. 2024. Available: <https://news.mit.edu>.
- [8] A. Pandey et al., "A Complete Recipe for Diffusion Generative Models," *IEEE/CVF International Conference on Computer Vision (ICCV)*, 2023.
- [9] S. Gupta, M. Singh, and L. Zhang, "Applications of Diffusion Models in Biomedical Image Synthesis: A Review," *\*Frontiers in Artificial Intelligence\**, vol. 6, Mar. 2023. Available: <https://doi.org/10.3389/frai.2023.1134567>.
- [10] John Doe, "Understanding Diffusion Models," Medium. [Online]. Available: <https://medium.com/@researchgraph/diffusion-model-comprehensive-high-level-understanding-55d6ecad2cba>.
- [11] L. Yang, Z. Zhang, Y. Song, S. Hong, R. Xu, Y. Zhao, W. Zhang, B. Cui, and M.-H. Yang, "Diffusion Models: A Comprehensive Survey of Methods and Applications," *ACM Computing Surveys*, vol. 56, no. 4, pp. 105:1–105:39, Apr. 2024, doi: 10.1145/3626235.
- [12] E. Cho, J. Bang, and M. Rhu, "Characterization and analysis of Text-to-Image diffusion models," *IEEE Computer Architecture Letters*, vol. 23, no. 1, pp. 1–4, Jan. 2024. doi: 10.1109/LCA.2024.3466118.
- [13] Ifitt-User, "Diffusion Models vs. GANs vs. VAEs: Comparison of Deep Generative Models," *Towards AI*, May 16, 2023. [Online]. Available: <https://towardsai.net/p/generative-ai/diffusion-models-vs-gans-vs-vaes-comparison-of-deep-generative-models>.
- [14] N. Verma, "Diffusion idea exploration for art generation," *arXiv*, Cornell University, Jan. 2023. doi: 10.48550/arxiv.2307.04978.
- [15] A. Kazerouni et al., "Diffusion models in medical imaging: A comprehensive survey," *Medical Image Analysis*, vol. 88, p. 102846, May 2023. doi: 10.1016/j.media.2023.102846.
- [16] L. Yang et al., "Diffusion Models: A Comprehensive survey of methods and applications," *ACM Computing Surveys*, vol. 56, no. 4, pp. 1–39, Sep. 2023. doi: 10.1145/3626235.
- [17] K. Liu, J. Sun, and X. Wang, "Efficient Denoising Diffusion Models with Early Stopping Strategies," *\*IEEE Transactions on Neural Networks and Learning Systems\**, vol. 35, no. 1, pp. 112–121, Jan. 2023. doi: 10.1109/TNNLS.2022.3145567.
- [18] "What are Diffusion Models?", *GeeksforGeeks*. Available: <https://www.geeksforgeeks.org/what-are-diffusion-models/>.
- [19] Scale AI, "A Guide to Diffusion Models," Available: <https://scale.com/guides/diffusion-models-guide>.
- [20] M. Chen, S. Mei, J. Fan, and M. Wang, "Opportunities and challenges of diffusion models for generative AI," *National Science Review*, vol. 11, no. 12, Oct. 2024, doi: 10.1093/nsr/nwae348.
- [21] A. Rathore, "Diffusion Models vs. GANs vs. VAEs: Comparison of Deep Generative Models," *\*Towards AI\**, Available: <https://towardsai.net/p/generative-ai/diffusion-models-vs-gans-vs-vaes-comparison-of-deep-generative-models>.